

exp-cnn-63130500042

November 28, 2023

```
[1]: from tensorflow.keras.utils import image_dataset_from_directory

import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A
NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy
(detected version 1.24.3
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
[2]: RAND_STATE_VALUE = 42
DATA_TEST_PATH = "/kaggle/input/landuse-scene-classification/
↳images_train_test_val/test"
DATA_TRAIN_PATH = "/kaggle/input/landuse-scene-classification/
↳images_train_test_val/train"
DATA_VAL_PATH = "/kaggle/input/landuse-scene-classification/
↳images_train_test_val/validation"
IMG_SIZE = (256, 256)
IMG_SHAPE = (256, 256, 3)
LABEL_MODE = "categorical"
```

1 Load Data

```
[3]: train_ds = image_dataset_from_directory(directory = DATA_TRAIN_PATH, label_mode=
↳ LABEL_MODE, image_size = IMG_SIZE)
test_ds = image_dataset_from_directory(directory = DATA_TEST_PATH, label_mode =
↳ LABEL_MODE, image_size = IMG_SIZE)
val_ds = image_dataset_from_directory(directory = DATA_VAL_PATH, label_mode =
↳ LABEL_MODE, image_size = IMG_SIZE)
```

```
Found 7350 files belonging to 21 classes.
Found 1050 files belonging to 21 classes.
Found 2100 files belonging to 21 classes.
```

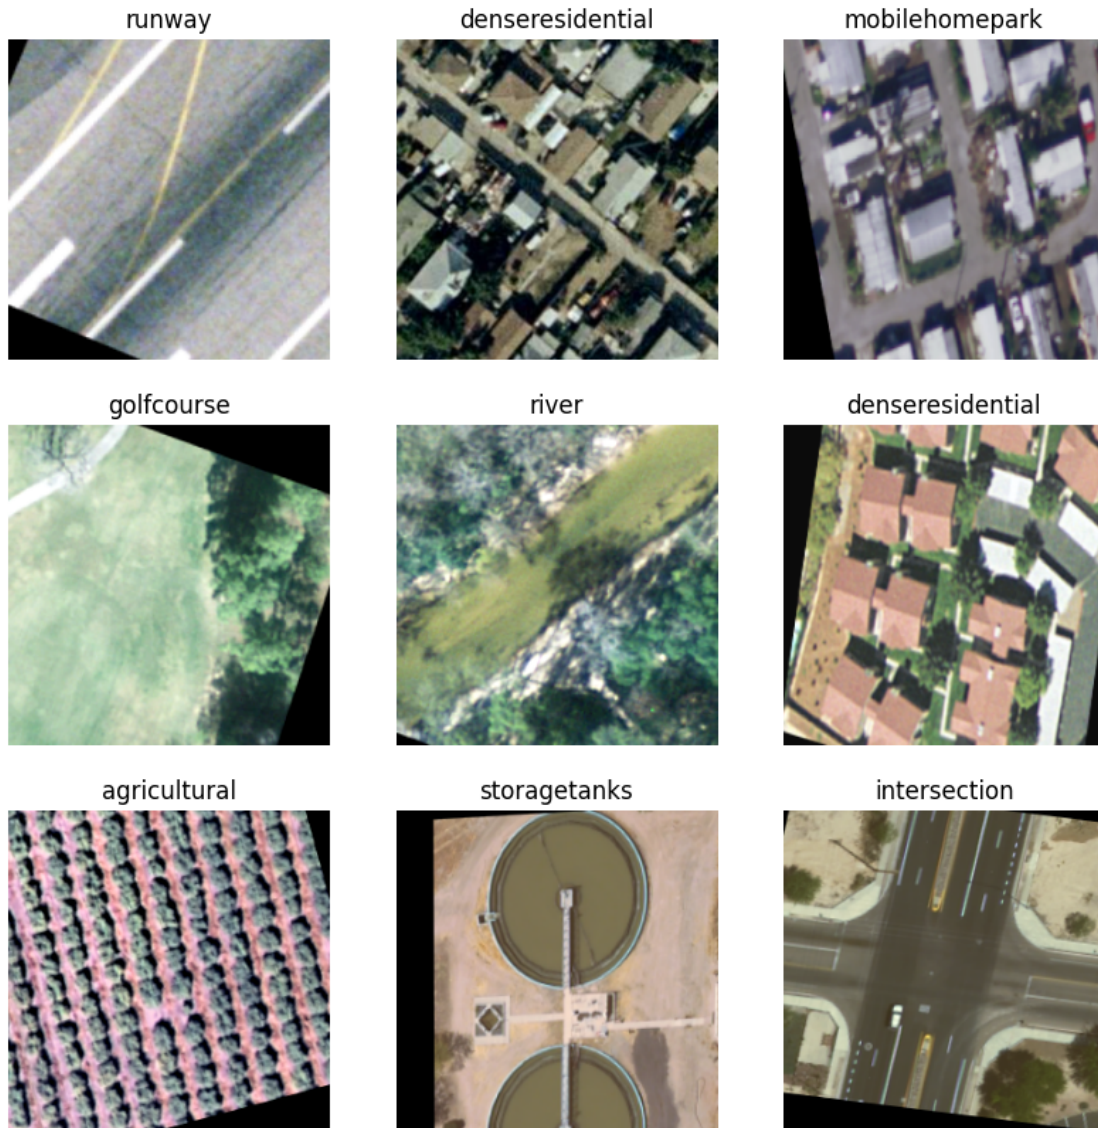
```
[4]: for image_batch, labels_batch in train_ds:
      print(image_batch.shape)
      print(labels_batch.shape)
      break
```

```
(32, 256, 256, 3)
(32, 21)
```

```
[5]: class_names = train_ds.class_names
      class_names
```

```
[5]: ['agricultural',
      'airplane',
      'baseballdiamond',
      'beach',
      'buildings',
      'chaparral',
      'denseresidential',
      'forest',
      'freeway',
      'golfcourse',
      'harbor',
      'intersection',
      'mediumresidential',
      'mobilehomepark',
      'overpass',
      'parkinglot',
      'river',
      'runway',
      'sparseresidential',
      'storagetanks',
      'tenniscourt']
```

```
[6]: plt.figure(figsize = (10, 10))
      for images, labels in train_ds.take(1):
          for i in range(9):
              ax = plt.subplot(3, 3, i + 1)
              plt.imshow(images[i].numpy().astype("uint8"))
              plt.title(class_names[np.argmax(labels[i])])
              plt.axis("off")
      plt.show()
```



2 Model Creation

```
[7]: from tensorflow.keras import Sequential
      from tensorflow.keras import layers
      from tensorflow.keras.regularizers import l2
```

```
[8]: data_augmentation = Sequential([
      layers.RandomFlip("horizontal", input_shape = IMG_SHAPE),
      layers.RandomRotation(0.1),
      layers.RandomZoom(0.1)
    ])
```

```
[9]: model = Sequential([
    data_augmentation,
    layers.Rescaling(1./255),
    layers.Conv2D(32, (3, 3), activation = "relu", kernel_regularizer = l2(0.
↪0001)),
    layers.MaxPooling2D(),
    layers.BatchNormalization(),
    layers.Conv2D(64, (3, 3), activation = "relu", kernel_regularizer = l2(0.
↪0001)),
    layers.MaxPooling2D(),
    layers.BatchNormalization(),
    layers.Conv2D(128, (3, 3), activation = "relu", kernel_regularizer = l2(0.
↪0001)),
    layers.MaxPooling2D(),
    layers.BatchNormalization(),
    layers.Dropout(0.25),
    layers.Flatten(),
    layers.Dense(256, activation = "relu"),
    layers.Dense(len(train_ds.class_names), activation = "softmax")
])
```

```
[10]: model.compile(optimizer = "adam", loss = "categorical_crossentropy", metrics =_
↪["accuracy"])
```

```
[11]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 256, 256, 3)	0
rescaling (Rescaling)	(None, 256, 256, 3)	0
conv2d (Conv2D)	(None, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
batch_normalization (Batch Normalization)	(None, 127, 127, 32)	128
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0

batch_normalization_1 (Batch Normalization)	(None, 62, 62, 64)	256
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
batch_normalization_2 (Batch Normalization)	(None, 30, 30, 128)	512
dropout (Dropout)	(None, 30, 30, 128)	0
flatten (Flatten)	(None, 115200)	0
dense (Dense)	(None, 256)	29491456
dense_1 (Dense)	(None, 21)	5397

```

=====
Total params: 29590997 (112.88 MB)
Trainable params: 29590549 (112.88 MB)
Non-trainable params: 448 (1.75 KB)
-----

```

```
[12]: history = model.fit(train_ds, epochs = 100, validation_data = val_ds, verbose = 1)
```

Epoch 1/100

```

2023-11-27 22:17:29.409536: E
tensorflow/core/grappler/optimizers/meta_optimizer.cc:954] layout failed:
INVALID_ARGUMENT: Size of values 0 does not match size of permutation 4 @ fanin
shape insequential_1/dropout/dropout/SelectV2-2-TransposeNHWCtoNCHW-
LayoutOptimizer

```

```

230/230 [=====] - 26s 86ms/step - loss: 16.5802 -
accuracy: 0.1419 - val_loss: 13.3509 - val_accuracy: 0.0571

```

Epoch 2/100

```

230/230 [=====] - 19s 80ms/step - loss: 4.9771 -
accuracy: 0.1444 - val_loss: 3.3783 - val_accuracy: 0.1286

```

Epoch 3/100

```

230/230 [=====] - 19s 80ms/step - loss: 3.1955 -
accuracy: 0.1565 - val_loss: 2.7758 - val_accuracy: 0.1471

```

Epoch 4/100

```

230/230 [=====] - 19s 81ms/step - loss: 2.7289 -
accuracy: 0.2033 - val_loss: 3.1069 - val_accuracy: 0.2033

```

Epoch 5/100

```

230/230 [=====] - 19s 80ms/step - loss: 2.5577 -

```

accuracy: 0.2135 - val_loss: 2.8194 - val_accuracy: 0.2167
 Epoch 6/100
 230/230 [=====] - 19s 81ms/step - loss: 2.5004 -
 accuracy: 0.2297 - val_loss: 4.0873 - val_accuracy: 0.1810
 Epoch 7/100
 230/230 [=====] - 19s 81ms/step - loss: 2.4138 -
 accuracy: 0.2427 - val_loss: 2.6495 - val_accuracy: 0.2748
 Epoch 8/100
 230/230 [=====] - 19s 81ms/step - loss: 2.3582 -
 accuracy: 0.2620 - val_loss: 2.4668 - val_accuracy: 0.2067
 Epoch 9/100
 230/230 [=====] - 19s 82ms/step - loss: 2.3044 -
 accuracy: 0.2683 - val_loss: 2.5779 - val_accuracy: 0.2243
 Epoch 10/100
 230/230 [=====] - 19s 80ms/step - loss: 2.3033 -
 accuracy: 0.2770 - val_loss: 3.0955 - val_accuracy: 0.1943
 Epoch 11/100
 230/230 [=====] - 19s 81ms/step - loss: 2.2629 -
 accuracy: 0.2882 - val_loss: 2.6829 - val_accuracy: 0.2052
 Epoch 12/100
 230/230 [=====] - 19s 80ms/step - loss: 2.2153 -
 accuracy: 0.3014 - val_loss: 2.4648 - val_accuracy: 0.2657
 Epoch 13/100
 230/230 [=====] - 19s 80ms/step - loss: 2.1489 -
 accuracy: 0.3190 - val_loss: 2.2258 - val_accuracy: 0.3071
 Epoch 14/100
 230/230 [=====] - 19s 81ms/step - loss: 2.1002 -
 accuracy: 0.3328 - val_loss: 2.4760 - val_accuracy: 0.3052
 Epoch 15/100
 230/230 [=====] - 19s 80ms/step - loss: 2.1073 -
 accuracy: 0.3407 - val_loss: 3.7626 - val_accuracy: 0.2619
 Epoch 16/100
 230/230 [=====] - 19s 80ms/step - loss: 2.0555 -
 accuracy: 0.3457 - val_loss: 2.2917 - val_accuracy: 0.3367
 Epoch 17/100
 230/230 [=====] - 19s 80ms/step - loss: 2.0457 -
 accuracy: 0.3584 - val_loss: 2.0553 - val_accuracy: 0.3476
 Epoch 18/100
 230/230 [=====] - 19s 80ms/step - loss: 1.9973 -
 accuracy: 0.3686 - val_loss: 2.0705 - val_accuracy: 0.3667
 Epoch 19/100
 230/230 [=====] - 19s 81ms/step - loss: 1.9695 -
 accuracy: 0.3737 - val_loss: 2.2393 - val_accuracy: 0.3619
 Epoch 20/100
 230/230 [=====] - 19s 80ms/step - loss: 1.9480 -
 accuracy: 0.3795 - val_loss: 1.9986 - val_accuracy: 0.3652
 Epoch 21/100
 230/230 [=====] - 19s 81ms/step - loss: 1.9281 -

accuracy: 0.3893 - val_loss: 3.4114 - val_accuracy: 0.2814
 Epoch 22/100
 230/230 [=====] - 19s 80ms/step - loss: 1.9177 -
 accuracy: 0.4018 - val_loss: 1.9023 - val_accuracy: 0.4081
 Epoch 23/100
 230/230 [=====] - 19s 80ms/step - loss: 1.8712 -
 accuracy: 0.4177 - val_loss: 2.1609 - val_accuracy: 0.3719
 Epoch 24/100
 230/230 [=====] - 19s 81ms/step - loss: 1.8374 -
 accuracy: 0.4186 - val_loss: 1.8434 - val_accuracy: 0.4295
 Epoch 25/100
 230/230 [=====] - 19s 80ms/step - loss: 1.8069 -
 accuracy: 0.4246 - val_loss: 2.5232 - val_accuracy: 0.3067
 Epoch 26/100
 230/230 [=====] - 19s 80ms/step - loss: 1.8157 -
 accuracy: 0.4244 - val_loss: 1.9525 - val_accuracy: 0.4229
 Epoch 27/100
 230/230 [=====] - 19s 81ms/step - loss: 1.7743 -
 accuracy: 0.4371 - val_loss: 1.9757 - val_accuracy: 0.4390
 Epoch 28/100
 230/230 [=====] - 19s 80ms/step - loss: 1.7348 -
 accuracy: 0.4537 - val_loss: 2.4870 - val_accuracy: 0.3729
 Epoch 29/100
 230/230 [=====] - 19s 80ms/step - loss: 1.7664 -
 accuracy: 0.4437 - val_loss: 1.9859 - val_accuracy: 0.4090
 Epoch 30/100
 230/230 [=====] - 19s 80ms/step - loss: 1.7457 -
 accuracy: 0.4521 - val_loss: 1.8777 - val_accuracy: 0.4386
 Epoch 31/100
 230/230 [=====] - 19s 81ms/step - loss: 1.6833 -
 accuracy: 0.4735 - val_loss: 1.8890 - val_accuracy: 0.4605
 Epoch 32/100
 230/230 [=====] - 19s 80ms/step - loss: 1.6808 -
 accuracy: 0.4720 - val_loss: 1.9725 - val_accuracy: 0.4286
 Epoch 33/100
 230/230 [=====] - 19s 80ms/step - loss: 1.6810 -
 accuracy: 0.4799 - val_loss: 1.8703 - val_accuracy: 0.4405
 Epoch 34/100
 230/230 [=====] - 19s 80ms/step - loss: 1.6508 -
 accuracy: 0.4838 - val_loss: 1.8363 - val_accuracy: 0.4657
 Epoch 35/100
 230/230 [=====] - 19s 80ms/step - loss: 1.6338 -
 accuracy: 0.4936 - val_loss: 1.7103 - val_accuracy: 0.4900
 Epoch 36/100
 230/230 [=====] - 19s 81ms/step - loss: 1.5897 -
 accuracy: 0.5026 - val_loss: 2.1301 - val_accuracy: 0.4129
 Epoch 37/100
 230/230 [=====] - 19s 81ms/step - loss: 1.5871 -

accuracy: 0.5050 - val_loss: 2.0472 - val_accuracy: 0.4219
Epoch 38/100
230/230 [=====] - 19s 80ms/step - loss: 1.5517 -
accuracy: 0.5261 - val_loss: 1.7969 - val_accuracy: 0.4838
Epoch 39/100
230/230 [=====] - 19s 80ms/step - loss: 1.5688 -
accuracy: 0.5180 - val_loss: 1.7578 - val_accuracy: 0.5005
Epoch 40/100
230/230 [=====] - 19s 81ms/step - loss: 1.5079 -
accuracy: 0.5378 - val_loss: 1.8278 - val_accuracy: 0.4790
Epoch 41/100
230/230 [=====] - 19s 81ms/step - loss: 1.5087 -
accuracy: 0.5322 - val_loss: 1.9791 - val_accuracy: 0.4576
Epoch 42/100
230/230 [=====] - 19s 80ms/step - loss: 1.5097 -
accuracy: 0.5405 - val_loss: 1.8369 - val_accuracy: 0.4576
Epoch 43/100
230/230 [=====] - 19s 80ms/step - loss: 1.5324 -
accuracy: 0.5332 - val_loss: 2.0184 - val_accuracy: 0.4733
Epoch 44/100
230/230 [=====] - 19s 81ms/step - loss: 1.4948 -
accuracy: 0.5423 - val_loss: 2.3227 - val_accuracy: 0.3995
Epoch 45/100
230/230 [=====] - 19s 80ms/step - loss: 1.4702 -
accuracy: 0.5491 - val_loss: 2.0370 - val_accuracy: 0.4724
Epoch 46/100
230/230 [=====] - 19s 81ms/step - loss: 1.4725 -
accuracy: 0.5548 - val_loss: 1.6260 - val_accuracy: 0.5424
Epoch 47/100
230/230 [=====] - 19s 80ms/step - loss: 1.4461 -
accuracy: 0.5604 - val_loss: 1.8315 - val_accuracy: 0.4867
Epoch 48/100
230/230 [=====] - 19s 81ms/step - loss: 1.4204 -
accuracy: 0.5676 - val_loss: 1.6233 - val_accuracy: 0.5281
Epoch 49/100
230/230 [=====] - 19s 80ms/step - loss: 1.4462 -
accuracy: 0.5599 - val_loss: 1.7576 - val_accuracy: 0.5152
Epoch 50/100
230/230 [=====] - 19s 81ms/step - loss: 1.4149 -
accuracy: 0.5737 - val_loss: 1.6091 - val_accuracy: 0.5290
Epoch 51/100
230/230 [=====] - 19s 81ms/step - loss: 1.3780 -
accuracy: 0.5810 - val_loss: 1.6019 - val_accuracy: 0.5590
Epoch 52/100
230/230 [=====] - 19s 82ms/step - loss: 1.3716 -
accuracy: 0.5837 - val_loss: 1.4176 - val_accuracy: 0.5814
Epoch 53/100
230/230 [=====] - 19s 81ms/step - loss: 1.3631 -


accuracy: 0.5895 - val_loss: 1.6844 - val_accuracy: 0.5162
 Epoch 54/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3702 -
 accuracy: 0.5853 - val_loss: 2.0403 - val_accuracy: 0.4971
 Epoch 55/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3677 -
 accuracy: 0.5867 - val_loss: 1.9503 - val_accuracy: 0.4719
 Epoch 56/100
 230/230 [=====] - 19s 81ms/step - loss: 1.3628 -
 accuracy: 0.5856 - val_loss: 1.9504 - val_accuracy: 0.4962
 Epoch 57/100
 230/230 [=====] - 19s 81ms/step - loss: 1.3249 -
 accuracy: 0.5999 - val_loss: 1.5258 - val_accuracy: 0.5771
 Epoch 58/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3354 -
 accuracy: 0.6027 - val_loss: 1.6401 - val_accuracy: 0.5395
 Epoch 59/100
 230/230 [=====] - 19s 81ms/step - loss: 1.3301 -
 accuracy: 0.6029 - val_loss: 1.4793 - val_accuracy: 0.5786
 Epoch 60/100
 230/230 [=====] - 19s 81ms/step - loss: 1.3041 -
 accuracy: 0.6084 - val_loss: 1.7342 - val_accuracy: 0.5133
 Epoch 61/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3197 -
 accuracy: 0.6073 - val_loss: 1.5512 - val_accuracy: 0.5490
 Epoch 62/100
 230/230 [=====] - 19s 81ms/step - loss: 1.3130 -
 accuracy: 0.6020 - val_loss: 1.8412 - val_accuracy: 0.5205
 Epoch 63/100
 230/230 [=====] - 19s 81ms/step - loss: 1.2990 -
 accuracy: 0.6086 - val_loss: 1.8606 - val_accuracy: 0.5233
 Epoch 64/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3048 -
 accuracy: 0.6087 - val_loss: 1.5594 - val_accuracy: 0.5652
 Epoch 65/100
 230/230 [=====] - 19s 80ms/step - loss: 1.3411 -
 accuracy: 0.6042 - val_loss: 1.6015 - val_accuracy: 0.5514
 Epoch 66/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2824 -
 accuracy: 0.6205 - val_loss: 1.4710 - val_accuracy: 0.5786
 Epoch 67/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2872 -
 accuracy: 0.6110 - val_loss: 1.5829 - val_accuracy: 0.5495
 Epoch 68/100
 230/230 [=====] - 19s 82ms/step - loss: 1.2675 -
 accuracy: 0.6190 - val_loss: 1.4142 - val_accuracy: 0.5957
 Epoch 69/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2748 -

accuracy: 0.6216 - val_loss: 2.4479 - val_accuracy: 0.5057
Epoch 70/100
230/230 [=====] - 19s 80ms/step - loss: 1.2555 -
accuracy: 0.6223 - val_loss: 1.5017 - val_accuracy: 0.5733
Epoch 71/100
230/230 [=====] - 19s 80ms/step - loss: 1.2776 -
accuracy: 0.6219 - val_loss: 1.5575 - val_accuracy: 0.5705
Epoch 72/100
230/230 [=====] - 19s 81ms/step - loss: 1.2611 -
accuracy: 0.6291 - val_loss: 1.5925 - val_accuracy: 0.5681
Epoch 73/100
230/230 [=====] - 19s 82ms/step - loss: 1.2529 -
accuracy: 0.6294 - val_loss: 1.6489 - val_accuracy: 0.5743
Epoch 74/100
230/230 [=====] - 19s 80ms/step - loss: 1.2491 -
accuracy: 0.6303 - val_loss: 1.4689 - val_accuracy: 0.5910
Epoch 75/100
230/230 [=====] - 19s 80ms/step - loss: 1.2379 -
accuracy: 0.6348 - val_loss: 1.4342 - val_accuracy: 0.6105
Epoch 76/100
230/230 [=====] - 19s 81ms/step - loss: 1.2254 -
accuracy: 0.6427 - val_loss: 1.4764 - val_accuracy: 0.5824
Epoch 77/100
230/230 [=====] - 19s 80ms/step - loss: 1.2148 -
accuracy: 0.6422 - val_loss: 1.4682 - val_accuracy: 0.6010
Epoch 78/100
230/230 [=====] - 19s 82ms/step - loss: 1.2293 -
accuracy: 0.6314 - val_loss: 1.4897 - val_accuracy: 0.5952
Epoch 79/100
230/230 [=====] - 19s 81ms/step - loss: 1.2380 -
accuracy: 0.6348 - val_loss: 1.3853 - val_accuracy: 0.6062
Epoch 80/100
230/230 [=====] - 19s 80ms/step - loss: 1.2240 -
accuracy: 0.6393 - val_loss: 1.5829 - val_accuracy: 0.5757
Epoch 81/100
230/230 [=====] - 19s 80ms/step - loss: 1.2363 -
accuracy: 0.6415 - val_loss: 1.3228 - val_accuracy: 0.6100
Epoch 82/100
230/230 [=====] - 19s 81ms/step - loss: 1.2218 -
accuracy: 0.6446 - val_loss: 1.5766 - val_accuracy: 0.5690
Epoch 83/100
230/230 [=====] - 19s 81ms/step - loss: 1.2140 -
accuracy: 0.6502 - val_loss: 1.4977 - val_accuracy: 0.5933
Epoch 84/100
230/230 [=====] - 19s 80ms/step - loss: 1.2050 -
accuracy: 0.6403 - val_loss: 1.5543 - val_accuracy: 0.5867
Epoch 85/100
230/230 [=====] - 19s 80ms/step - loss: 1.1916 -

accuracy: 0.6465 - val_loss: 1.4223 - val_accuracy: 0.6052
 Epoch 86/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2145 -
 accuracy: 0.6510 - val_loss: 1.3968 - val_accuracy: 0.6176
 Epoch 87/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2229 -
 accuracy: 0.6449 - val_loss: 1.5621 - val_accuracy: 0.5810
 Epoch 88/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1907 -
 accuracy: 0.6567 - val_loss: 1.5821 - val_accuracy: 0.5900
 Epoch 89/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2149 -
 accuracy: 0.6521 - val_loss: 1.6131 - val_accuracy: 0.5652
 Epoch 90/100
 230/230 [=====] - 19s 80ms/step - loss: 1.2228 -
 accuracy: 0.6416 - val_loss: 1.3954 - val_accuracy: 0.6143
 Epoch 91/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1908 -
 accuracy: 0.6498 - val_loss: 1.3755 - val_accuracy: 0.6195
 Epoch 92/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1762 -
 accuracy: 0.6604 - val_loss: 1.3655 - val_accuracy: 0.6129
 Epoch 93/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1612 -
 accuracy: 0.6631 - val_loss: 1.5934 - val_accuracy: 0.5633
 Epoch 94/100
 230/230 [=====] - 19s 81ms/step - loss: 1.1908 -
 accuracy: 0.6547 - val_loss: 1.4892 - val_accuracy: 0.6062
 Epoch 95/100
 230/230 [=====] - 19s 81ms/step - loss: 1.1781 -
 accuracy: 0.6578 - val_loss: 1.3628 - val_accuracy: 0.6262
 Epoch 96/100
 230/230 [=====] - 19s 81ms/step - loss: 1.1916 -
 accuracy: 0.6520 - val_loss: 1.3125 - val_accuracy: 0.6276
 Epoch 97/100
 230/230 [=====] - 19s 81ms/step - loss: 1.1753 -
 accuracy: 0.6630 - val_loss: 1.5007 - val_accuracy: 0.5924
 Epoch 98/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1821 -
 accuracy: 0.6562 - val_loss: 1.3931 - val_accuracy: 0.6181
 Epoch 99/100
 230/230 [=====] - 19s 80ms/step - loss: 1.1852 -
 accuracy: 0.6565 - val_loss: 1.3974 - val_accuracy: 0.6176
 Epoch 100/100
 230/230 [=====] - 19s 81ms/step - loss: 1.1713 -
 accuracy: 0.6634 - val_loss: 1.7742 - val_accuracy: 0.5643

```
[13]: # Plot training history for accuracy
plt.subplot(2, 1, 1)
plt.plot(history.history["accuracy"], label = "Training Accuracy")
plt.plot(history.history["val_accuracy"], label = "Validation Accuracy")
plt.title("Training and Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.legend()

# Plot training history for loss
plt.subplot(2, 1, 2)
plt.plot(history.history["loss"], label = "Training Loss")
plt.plot(history.history["val_loss"], label = "Validation Loss")
plt.title("Training and Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.tight_layout()
plt.show()
```



3 Model Testing and Evaluation

```
[14]: from sklearn.metrics import confusion_matrix
      from sklearn.metrics import classification_report
```

```
[ ]: def decode_labels(labels):
      return np.argmax(labels, axis = 1)
```

```
[ ]: true_labels = []
      all_pred_classes = []

      for images, labels in test_ds:
          predictions = model.predict(images)
          pred_classes = decode_labels(predictions)

          true_labels.extend(decode_labels(labels))
          all_pred_classes.extend(pred_classes)
```

```
1/1 [=====] - 0s 169ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
```

```

1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 342ms/step


```

```

[ ]: cmt = confusion_matrix(true_labels, all_pred_classes)

plt.figure(figsize = (10, 8))
sns.heatmap(cmt, annot = True, fmt = "d", cmap = "Blues", xticklabels = _
    ↪class_names, yticklabels = class_names)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

```



```
[ ]: report = classification_report(true_labels, all_pred_classes, target_names =  
    ↪class_names)  
print("Classification Report:")  
print(report)
```

Classification Report:


	precision	recall	f1-score	support
agricultural	0.51	0.94	0.66	50
airplane	0.61	0.34	0.44	50
baseballdiamond	0.68	0.30	0.42	50
beach	0.89	0.64	0.74	50
buildings	0.48	0.32	0.39	50
chaparral	1.00	0.88	0.94	50
denseresidential	0.36	0.42	0.39	50
forest	0.55	1.00	0.71	50
freeway	0.61	0.76	0.68	50
golfcourse	0.45	0.34	0.39	50
harbor	0.96	0.90	0.93	50
intersection	0.75	0.76	0.75	50
mediumresidential	0.29	0.24	0.26	50
mobilehomepark	0.60	0.60	0.60	50
overpass	0.72	0.66	0.69	50
parkinglot	0.82	0.84	0.83	50
river	0.60	0.72	0.65	50
runway	0.43	0.42	0.42	50
sparseresidential	0.50	0.52	0.51	50
storagetanks	0.30	0.22	0.25	50
tenniscourt	0.37	0.46	0.41	50
accuracy			0.58	1050
macro avg	0.59	0.58	0.57	1050
weighted avg	0.59	0.58	0.57	1050

```
[21]: for images, labels in test_ds.take(1):  
    predictions = model.predict(images)  
    pred_classes = decode_labels(predictions)  
  
    plt.figure(figsize = (10, 10))  
    for i in range(9):  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        pred_class = class_names[pred_classes[i]]  
        true_class = class_names[np.argmax(labels[i])]  
        prob = np.max(predictions[i])  
        plt.title(f"Predict: {pred_class} ({prob:.2f})\n Actual: {true_class}")
```


```
plt.axis("off")  
plt.show()
```

1/1 [=====] - 0s 45ms/step

Predict: airplane (0.66)
Actual: buildings




Predict: agricultural (1.00)
Actual: agricultural



Predict: beach (0.67)
Actual: beach



Predict: agricultural (0.93)
Actual: beach



Predict: buildings (0.46)
Actual: airplane



Predict: intersection (0.60)
Actual: intersection



Predict: overpass (0.31)
Actual: storagetanks



Predict: sparseresidential (1.00)
Actual: sparseresidential



Predict: freeway (0.92)
Actual: freeway

